

Description of Facilities and Resources

Livermore Computing | Lawrence Livermore National Laboratory (LLNL)



LLNL computational scientists are supported by Livermore Computing (LC), which delivers a balanced High Performance Computing (HPC) environment with constantly evolving hardware resources and a wealth of HPC expertise in porting, running, and tuning real-world, large-scale applications. Currently LC delivers over 186 petaflops of compute power, massive shared parallel file systems, powerful data analysis platforms, and archival storage capable of storing hundreds of petabytes of data. This balanced hardware environment supports key collaborations between LLNL applications developers and LC experts. This involves the creation, production use, and performance monitoring and analysis of results of HPC parallel applications in a wide variety of scientific disciplines.

In addition to general access for development and debugging, researchers access allocations on these resources through several competitive internal programs. With these allocations, science and technology directorates sustain and enhance institutional capabilities to help ensure mission successes. Such efforts benefit the entire Laboratory by developing the expertise of the next generation of scientists and engineers who will, in turn, use these capabilities for continued growth of Laboratory programs. Furthermore, researchers with “work-for-others” (WFO) projects and external collaborators needing HPC access can purchase cycles on the Multi-programmatic and Institutional Computing systems.

Major systems include the heterogeneous 125-petaflops Sierra and its 23-petaflops unclassified sister system, Lassen, with a combined 164PB file system; Corona—which was increased to 11.5 petaflops to join the fight against COVID-19; the nearly 6-petaflops Ruby; the 5-petaflops Magma; the Jade and Quartz systems at 3+ petaflops each; and additional large multi-core, multisocket Linux clusters with a variety of processor types, ranging from IBM PowerPC to Intel Broadwell processors. Our El Capitan testbeds feature HPE/Cray Slingshot interconnect technology and contain AMD Epyc Rome processors and AMD MI-60 and MI-100 GPU accelerators. We also have two cutting-edge AI powerhouses: the SambaNova accelerator for Corona and the Cerebras wafer-scale engine integrated into Lassen. There are 22,530 NVIDIA Graphical Processing Units (GPUs) across several of the platforms. In total, more than 16,312 nodes, 695,752 cores, and 3.68 petabytes of memory are available across the production LC systems on two networks.

The balanced LLNL simulation environment includes Lustre multi-cluster file systems, HPSS-based archival resources, a high-speed Ethernet networking infrastructure, advanced visualization resources, and a rich tool environment, as described on page two.



PRIMARY PRODUCTION COMPUTING PLATFORMS

Sierra and Lassen

Sierra, with its 125 petaflops peak, combines two types of processor chip—IBM's Power 9 processors and NVIDIA's Volta GPUs. Sierra is designed for more efficient overall operations and is a leading architecture for extreme-scale computing. Lassen shares Sierra's unique architecture and is an unclassified system of 23 petaflops. Their smaller 1.5-petaflops sibling system is RZAnsel.

Corona

An 11.5 petaflops heterogeneous Penguin Computing system comprised of AMD processors and AMD GPU accelerators. Corona has been upgraded for the fight against the novel coronavirus.

Ruby

A 6-petaflops CPU-only Supermicro cluster with a Cornelis-Networks OmniPath interconnect and 290 terabytes of memory.

Magma

A 5.3 petaflops Penguin Computing "Relion" system comprised of 752 nodes with Intel Xeon Platinum 9242 processors. The cluster has 293 terabytes of memory, liquid cooling, and an Intel Omni-Path interconnect.

Quartz

Quartz, Penguin system, has 3,018 nodes, 344 terabytes of memory and a peak speed of 3.7 petaflops

Jade

Jade, a Penguin Computing system, has 2,688 nodes and 343 terabytes of memory and a peak speed of 3.2 petaflops. It has Intel Broadwell processors. RZTopaz is a smaller 748-node cluster with the same architecture as Jade.

Syrah

This Cray system is a multi-core, multi-socket Linux cluster. Syrah has 324 nodes and 20 terabytes of memory, clocking in at a peak speed of 108 teraflops.

For a complete list of more than 20 production compute platforms supported by LC, see the [Livermore Computing Systems Summary](#).

INSTITUTIONAL RESOURCES AND FACILITIES

LLNL researchers benefit from an institutional IT infrastructure that provides desktop support and experts in server technologies. The latter includes virtualization expertise that has been applied to provide multiple operating systems on shared resources and to create a wide variety of virtual machines to leverage resources servers across LLNL. Networking service is also provided by an enterprise team. Connections into LLNL include ESnet and a wide variety of programmatic networks.

LC RESOURCES

High Performance Archival Storage

LC provides high-performance archival storage services via High Performance Storage System (HPSS). A world-class array of hardware integrated beneath HPSS includes disk arrays, tape subsystems, mover nodes, networks, robotics, and hundreds of petabytes of media.

File Systems

HPC systems require high performance storage, which LC provides in multiple forms: open-source Lustre, IBM Spectrum Scale, and the all-flash VAST. These systems are mounted across multiple compute clusters and deliver high-performance, global access to data.

LC Networking

LC's simulation environment includes a software-defined 100G high-speed datacenter Ethernet networking infrastructure as well as InfiniBand SANs. LC employs InfiniBand and Omni- Path fabrics for high-speed interconnects. Testbed work includes evaluation of next generation networking equipment, including the HPE/Cray Slingshot interconnect.

Containerization

LC has deployed a selection of on-prem-cloud science-focused platform capabilities to enable persistent data services, HPC workflow tools, and machine learning frameworks. These services will include container management and orchestration and an object storage system.

Visualization facilities

LC runs two large visualization clusters—the CPU-only Tron on the SCF with a peak speed of 434 teraflops and the hybrid-architecture Pascal on the OCF with a peak speed of 1.9 petaflops—as well as visualization partitions on both Sierra and Lassen. LC supports RealVNC on login nodes and NICE DCV on batch nodes and operates several visualization theaters, ranging from auditoriums with PowerWalls to smaller collaboration spaces.